

# Machine Intelligence & Deep Learning Workshop

Raymond Ptucha, Majid Rabbani, Mark Smith

The Kate Gleason **COLLEGE OF  
ENGINEERING**

## Recurrent Neural Networks



Raymond Ptucha  
June 27-29, 2018

Rochester Institute of Technology  
[www.rit.edu/kgcoe/cqas/machinelearning](http://www.rit.edu/kgcoe/cqas/machinelearning)



© 2018 Ray Ptucha, Rochester Institute of Technology

1

## Fair Use Agreement

This agreement covers the use of all slides in this document, please read carefully.

- You may freely use these slides for personal use, if:
  - My name (R. Ptucha) appears on each slide.
- You may freely use these slides externally, if:
  - You send me an email telling me the conference/venue/company name in advance, and which slides you wish to use.
  - You receive a positive confirmation email back from me.
  - My name (R. Ptucha) appears on each slide you use.

(c) Raymond Ptucha, [rwpeec@rit.edu](mailto:rwpeec@rit.edu)

© 2018 Ray Ptucha, Rochester Institute of Technology

2

# Agenda

- **Wed, June 27**
  - 9-10:30am Regression and Classification
  - 10:30-10:45pm Break
  - 10:45-12:15pm Boosting and SVM
  - 12:15-1:30pm Lunch
  - 1:30-3:30pm Neural Networks and Dimensionality Reduction
  - 3:30-5pm Hands-on Python and Machine Learning
- **Thur, June 28**
  - 9-10:30am Introduction to deep learning
  - 10:30-10:45pm Break
  - 10:45-12:15pm Convolutional Neural Networks
  - 12:15-1:30pm Lunch
  - 1:30-3:30pm Region and pixel-level convolutions
  - 3:30-5pm Hands-on CNNs
- **Fri, June 29**
  - 9-10:30am **Recurrent neural networks**
  - 10:30-10:45pm Break
  - 10:45-12:15pm Language and Vision
  - 12:15-1:30pm Lunch
  - 1:30-3:30pm Graph convolutional neural networks; Generative adversarial networks
  - 3:30-5pm Hands-on regional CNNs, RNNs

© 2018 Ray Ptucha, Rochester Institute of Technology

3

## Two Most Important Deep Learning Fields

- **Convolutional Neural Networks (CNN)**
  - Examine high dimensional input, learn features and classifier simultaneously
- **Recurrent Neural Networks (RNN)**
  - Learn temporal signals, remember both short and long sequences

© 2018 Ray Ptucha, Rochester Institute of Technology

4


# Recurrent Neural Networks

- Feed forward Artificial Neural Networks (ANNs) are great at classification, but are limited at predicting future given the past.
- Need framework that determines output based upon current and previous inputs.
- Recurrent or Recursive Neural Networks (RNNs) capture sequential information and are used in speech recognition, activity recognition, NLP, weather prediction, etc.

© 2018 Ray Ptucha, Rochester Institute of Technology

5


## Adding Recurrence



© 2018 Ray Ptucha, Rochester Institute of Technology

6

# Recurrent Networks



$$\begin{aligned} in_{h0} &= (W_{xh}x_t \\ h_t &= f(in_{h0}) \end{aligned} )$$


Where:

- $f$  is some activation function
- $x_t$  and  $h_t$  are current input and current output values
- $W_{xh}$  is the weight matrix for input, hidden and output stages respectively
- $in_{h0}$  is the input to activation function in hidden and output layers

© 2018 Ray Ptucha, Rochester Institute of Technology

7

# Recurrent Networks



$$\begin{aligned} in_{h0} &= (W_{xh}x_t + W_{hh}h_{t-1}) \\ h_t &= f(in_{h0}) \\ in_{y0} &= W_{hy}h_t \\ y_t &= f(in_{y0}) \end{aligned}$$

Where:

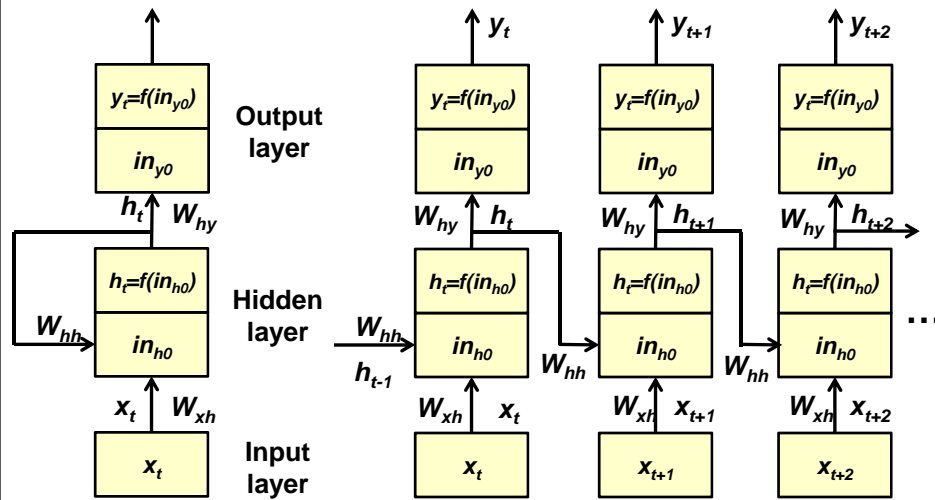
- $f$  is some activation function
- $x_t, h_t, h_{t-1}$  and  $y_t$  are current input, hidden, previous hidden and current output values
- $W_{xh}, W_{hh}$  and  $W_{hy}$  are the weight matrices for input, hidden and output stages respectively
- $in_{h0}$  and  $in_{y0}$  are the inputs to activation function in hidden and output layers

© 2018 Ray Ptucha, Rochester Institute of Technology

8

# Recurrent Networks

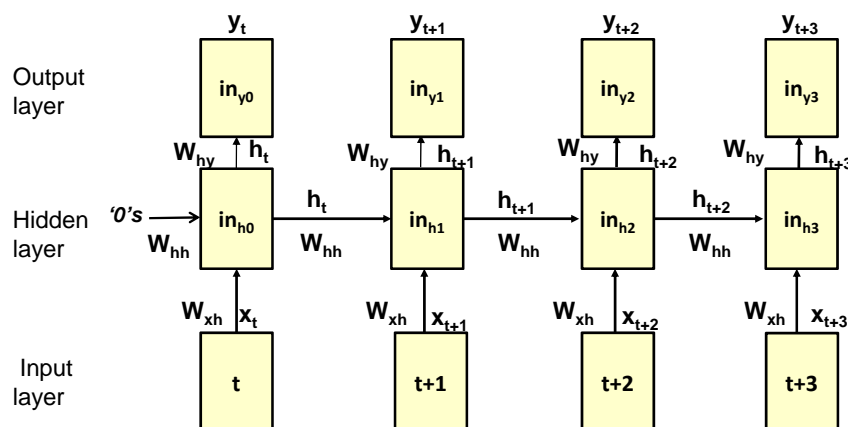
Both figures represent the same architecture



© 2018 Ray Ptucha, Rochester Institute of Technology

9

## Forward Propagation of Recurrent Networks




Note: regardless of how many time steps taken, only learning a single  $W_{xh}$ ,  $W_{hh}$ , and  $W_{hy}$ . Each are learned via standard back propagation.

© 2018 Ray Ptucha, Rochester Institute of Technology

10

# Recurrent Networks

Recurrent Neural Network "neuron"



$P(\text{next event} \mid \text{previous events})$


- Unfortunately, these vanilla RNNs don't always work.
- Can't store info over long periods of time.
- Suffer from vanishing and/or exploding gradients.

© 2018 Ray Ptucha, Rochester Institute of Technology


11

# Recurrent Networks

Recurrent Neural Network "neuron"



Long Short Term Memory "neuron"



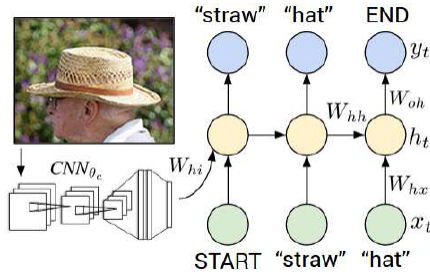
Donahue et al., 2015

- LSTM's allow read/write/reset functions to neurons.
- Remember past to predict the future- (over long time periods).
- Can have many hidden neurons per layer and many layers.

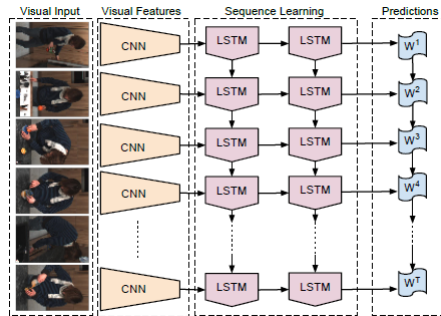
© 2018 Ray Ptucha, Rochester Institute of Technology

12

# Recurrent Applications

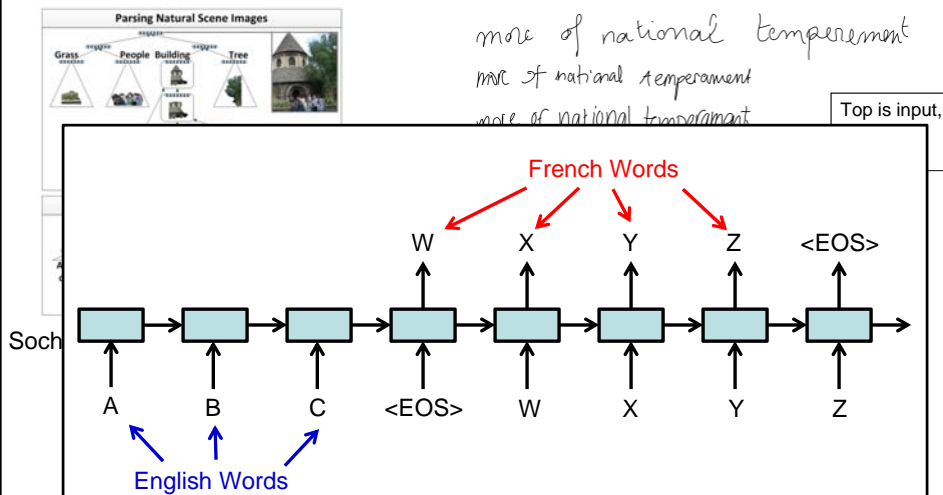


Karpathy, Fei-Fei, 2015




Donahue, et al., 2015

# Recurrent Applications



Sutskever et al., 2014

# Many Flavors




<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

© 2018 Ray Ptucha, Rochester Institute of Technology

15

# Vanilla Recurrent Neural Network



© 2018 Ray Ptucha, Rochester Institute of Technology



16



# Toy Example

- Character level language model
- Vocabulary: [h,e,l,o]
- Input training example “hello”

Embedding layer: For words, can use one-hot encoding, word2vec, glove, or solve via backprop




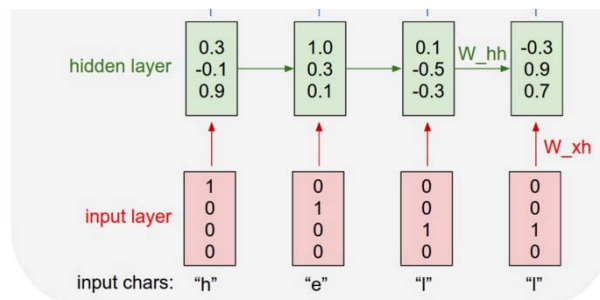
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

© 2018 Ray Ptucha, Rochester Institute of Technology

17

# Toy Example

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$



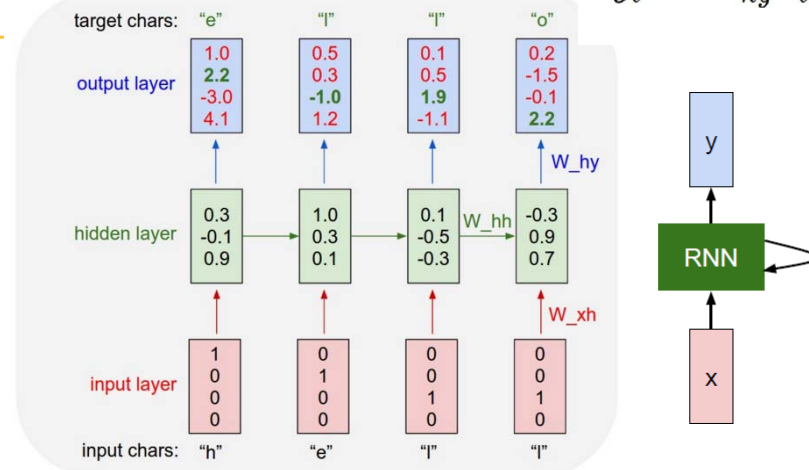
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

© 2018 Ray Ptucha, Rochester Institute of Technology

18

## Toy Example

$$y_t = W_{hy}h_t$$

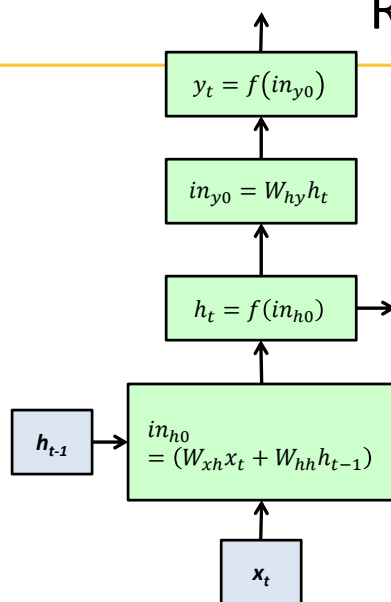


Example RNN with 4-dimensional input and output layers, and a hidden layer of 3 units (neurons). Shown are the activations in the forward pass when the RNN is fed the characters "hell" as input. The output layer contains confidences the RNN assigns for the next character (vocabulary is "h,e,l,o"); want the green numbers to be high and red numbers to be low <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

© 2018 Ray Ptucha, Rochester Institute of Technology

19

## Recurrent Networks



$$in_{h0} = (W_{xh}x_t + W_{hh}h_{t-1})$$

$$h_t = f(in_{h0})$$

$$in_{y0} = W_{hy}h_t$$

$$y_t = f(in_{y0})$$


Where:

- $f$  is some activation function
- $x_t, h_t, h_{t-1}$  and  $y_t$  are current input, hidden, previous hidden and current output values
- $W_{xh}, W_{hh}$  and  $W_{hy}$  are the weight matrices for input, hidden and output stages respectively
- $in_{h0}$  and  $in_{y0}$  are the inputs to activation function in hidden and output layers

© 2018 Ray Ptucha, Rochester Institute of Technology

22

# Recurrent Networks for Character Prediction



© 2018 Ray Ptucha, Rochester Institute of Technology

23

Karpathy's minimal character-level language model with a vanilla recurrent neural network in python. This is his RNN in 112 lines of code, <https://gist.github.com/karpathy/d4dee566867f8291f086>

GitHub Gist Search... All gists GitHub

Instantly share code, notes, and snippets.

karpathy / min-char-rnn.py  
Last active 5 hours ago

Code Revisions 6 Stars 368 Forks 120 Embed <script src="htt

Minimal character-level language model with a Vanilla Recurrent Neural Network, in Python/numpy

```

1  """
2  Minimal character-level Vanilla RNN model. Written by Andrej Karpathy (@karpathy)
3  BSD License
4  """
5  import numpy as np
6
7  # data I/O
8  data = open('input.txt', 'r').read() # should be simple plain text file
9  chars = list(set(data))
10 data_size, vocab_size = len(data), len(chars)
11 print 'data has %d characters, %d unique.' % (data_size, vocab_size)
12 char_to_ix = { ch:i for i,ch in enumerate(chars) }
```

© 2018 Ray Ptucha, Rochester Institute of Technology

24

Karpathy's minimal character-level language model with a vanilla recurrent neural network in python. This is his RNN in 112 lines of code, <https://gist.github.com/karpathy/d4dee566867f8291f086>

- Reads in a training set (say a novel from your favorite author), then uses RNN's to learn character prediction.
- Given the past characters, what is the next most likely character.
- We will see, this can learn words, sentences, phrases, etc. which look like they were written by the author of the novel!
- Given a 't', 'h' is the next most likely char; given 'th', 'e' is the most likely char; given 'the', ' ' is the most likely char, ...

© 2018 Ray Ptucha, Rochester Institute of Technology

25


## Vanishing gradient problem

- If the product of weights and derivatives of the activation function is less than 1 the gradients vanish rapidly with the time, else they explode.
- The longer the time lag the greater the chance of gradients exploding or vanishing.
- However weights tend to be initialized with a mean value of 0 and standard deviation of 1.
- Derivatives of common activation functions like sigmoid and tanh also tend to have a value less than or equal to 1 in most cases.
- Hence vanishing gradient problem is observed more commonly.

© 2018 Ray Ptucha, Rochester Institute of Technology

27

# LSTMs

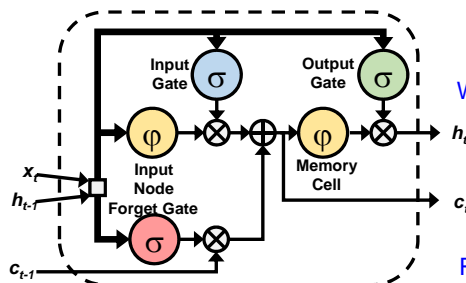


© 2018 Ray Ptucha, Rochester Institute of Technology

28

# LSTMs

Convert standard neuron into a complex memory cell



With  $\sigma()$ =sigmoid activation function and  $\phi()$ =tanh activation function,  $x_t$  and the previous cell output  $h_{t-1}$  calculate:

Write, read, reset governors:

$$\text{Input gate: } i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1})$$

$$\text{Output gate: } o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1})$$

$$\text{Forget gate: } f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1})$$

Real input to memory cell:

$$\text{Input node: } g_t = \phi(W_{xc}x_t + W_{hc}h_{t-1})$$

Looks just like our RNN cell!

Calculate a memory cell which is the summation of the previous memory cell, governed by the forget gate and the input and previous output governed by independent combinations of the same:

$$c_t = (f_t c_{t-1} + i_t g_t)$$


Calculate a new hidden state, governed by the output gate:

$$h_t = o_t \phi(c_t)$$

© 2018 Ray Ptucha, Rochester Institute of Technology

29

The input node summarizes the input and past output, which will be governed by the input gate.



With  $\sigma()$ =sigmoid activation function and  $\phi()$ =tanh activation function,  $x_t$  and the previous cell output  $h_{t-1}$  calculate:

$$\text{Input gate: } i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1})$$

$$\text{Output gate: } o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1})$$

$$\text{Forget gate: } f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1})$$

$$\text{Input node: } g_t = \phi(W_{xc}x_t + W_{hc}h_{t-1})$$

Calculate a memory cell which is the summation of the previous memory cell, governed by the forget gate and the input and previous output governed by independent combinations of the same:

$$c_t = (f_t c_{t-1} + i_t g_t)$$


Calculate a new hidden state, governed by the output gate:

$$h_t = o_t \phi(c_t)$$

© 2018 Ray Ptucha, Rochester Institute of Technology

30

Write: The input gate gives the provision to determine importance of current input and past hidden state.



With  $\sigma()$ =sigmoid activation function and  $\phi()$ =tanh activation function,  $x_t$  and the previous cell output  $h_{t-1}$  calculate:

$$\text{Input gate: } i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1})$$

$$\text{Output gate: } o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1})$$

$$\text{Forget gate: } f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1})$$

$$\text{Modulation gate: } g_t = \phi(W_{xc}x_t + W_{hc}h_{t-1})$$

Calculate a memory cell which is the summation of the previous memory cell, governed by the forget gate and the input and previous output governed by independent combinations of the same:

$$c_t = (f_t c_{t-1} + i_t g_t)$$


Calculate a new hidden state, governed by the output gate:

$$h_t = o_t \phi(c_t)$$

© 2018 Ray Ptucha, Rochester Institute of Technology

31

Read: The output gate determines what parts of the cell output are necessary for the next time step.



With  $\sigma()$ =sigmoid activation function and  $\phi()$ =tanh activation function,  $x_t$  and the previous cell output  $h_{t-1}$  calculate:

$$\text{Input gate: } i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1})$$

$$\text{Output gate: } o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1})$$

$$\text{Forget gate: } f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1})$$

$$\text{Modulation gate: } g_t = \phi(W_{xc}x_t + W_{hc}h_{t-1})$$

Calculate a memory cell which is the summation of the previous memory cell, governed by the forget gate and the input and previous output governed by independent combinations of the same:

$$c_t = (f_t c_{t-1} + i_t g_t)$$

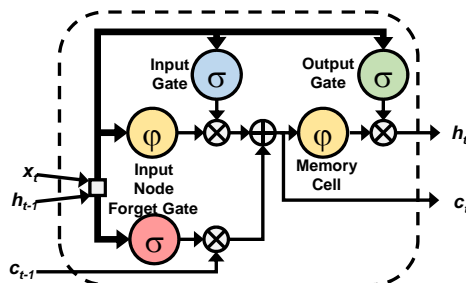
Calculate a new hidden state, governed by the output gate:

$$h_t = o_t \phi(c_t)$$

© 2018 Ray Ptucha, Rochester Institute of Technology

32

Reset: The forget gate gives the provision for the hidden layer to discard or forget the historical data



With  $\sigma()$ =sigmoid activation function and  $\phi()$ =tanh activation function,  $x_t$  and the previous cell output  $h_{t-1}$  calculate:

$$\text{Input gate: } i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1})$$

$$\text{Output gate: } o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1})$$

$$\text{Forget gate: } f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1})$$

$$\text{Modulation gate: } g_t = \phi(W_{xc}x_t + W_{hc}h_{t-1})$$

Calculate a memory cell which is the summation of the previous memory cell, governed by the forget gate and the input and previous output governed by independent combinations of the same:

$$c_t = (f_t c_{t-1} + i_t g_t)$$

Calculate a new hidden state, governed by the output gate:

$$h_t = o_t \phi(c_t)$$

© 2018 Ray Ptucha, Rochester Institute of Technology

33

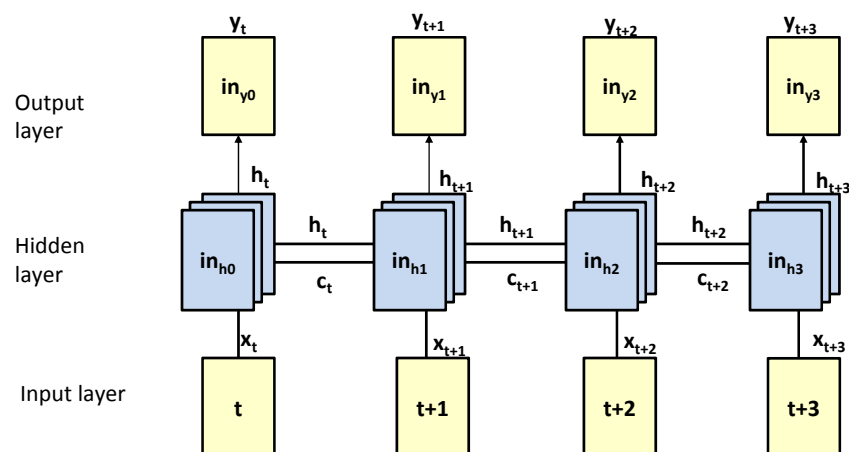
## Using LSTMs

- The LSTM memory cells are analogous to a single neuron.
- As such many hundreds of these memory cells are used in a layer, each of which passes its output  $h_t$  to the next time step,  $h_{t+1}$ .

© 2018 Ray Ptucha, Rochester Institute of Technology

34

Same architecture as RNNs, but middle neurons are now LSTM memory cells

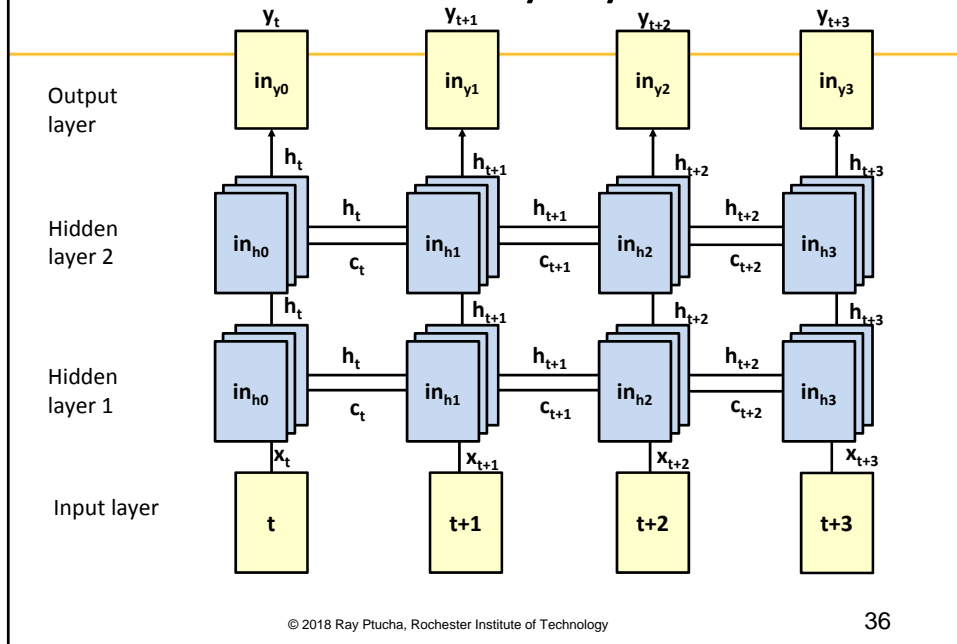


© 2018 Ray Ptucha, Rochester Institute of Technology

35




## Can do many layers...



36

## Advantages of LSTM over RNN

- As RNN cell step through time, the output of one cell is a transformed version of its input:  $h_t = f(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$
- With LSTMs, if we set the forget gate to 1:  $c_t = (f_t c_{t-1} + i_t g_t)$  the memory cell output can pass the entire previous input.
- Very similar to "skip layer" concept of ResNets



During backpropagation, all gradients forced to go through inverse gradient of  $f$ .

During backpropagation, gradients of  $c$  can skip  $f$ .


RNN gradients can vanish or explode, LSTM can still explode, so always use gradient clipping (restrict to  $\pm 5$ ).

© 2018 Ray Ptucha, Rochester Institute of Technology


37

# Many Variants of LSTMs

LSTM with CEC and input, output gates



Simplified variants: Gated Recurrent Unit



(b) Gated Recurrent Unit


K. Cho, B. van Merriënboer, D. Bahdanau, and Y. Bengio. On the properties of neural machine translation: Encoder-decoder approaches. arXiv preprint arXiv:1409.1259, 2014  
J. Chung, C. Gulcehre, K. Cho, Y. Bengio. Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling. arXiv preprint arXiv:1412.3555 (2014)

© 2018 Ray Ptucha, Rochester Institute of Technology

38

# Many Variants of LSTMs

LSTM: A Search Space Odyssey  
Greff et al., 2015



An Empirical Exploration of RNNs  
Jozefowicz et al., 2015

MUT1:

$$\begin{aligned} z &= \text{sign}(W_{xz}x_t + b_z) \\ r &= \text{sign}(W_{xr}x_t + W_{hr}h_t + b_r) \\ h_{t+1} &= \tanh(W_{hh}(r \odot h_t) + W_{zh}z_t + b_h) \odot z \\ &+ h_t \odot (1 - z) \end{aligned}$$

MUT2:

$$\begin{aligned} z &= \text{sign}(W_{xz}x_t + W_{hz}h_t + b_z) \\ r &= \text{sign}(r_x + W_{hr}h_t + b_r) \\ h_{t+1} &= \tanh(W_{hh}(r \odot h_t) + W_{zh}z_t + b_h) \odot z \\ &+ h_t \odot (1 - z) \end{aligned}$$

MUT3:

$$\begin{aligned} z &= \text{sign}(W_{xz}x_t + W_{hz} \tanh(h_t) + b_z) \\ r &= \text{sign}(W_{xr}x_t + W_{hr}h_t + b_r) \\ h_{t+1} &= \tanh(W_{hh}(r \odot h_t) + W_{zh}z_t + b_h) \odot z \\ &+ h_t \odot (1 - z) \end{aligned}$$

- Both tried many variants of LSTMs.
- Neither paper was able to find an architecture with substantial advantages over baseline Hochreiter et al. '97 formulation.

© 2018 Ray Ptucha, Rochester Institute of Technology

39

## GRU vs LSTM

Note: 1 update gate

LSTM uses separate update ( $i_t$ ) and forget ( $f_t$ ) gate values; then separate output  $h_t$  and memory states  $c_t$ .

### Gated Recurrent Unit

[Cho et al., EMNLP2014;  
Chung, Gulcehre, Cho, Bengio, DLUFL2014]

$$h_t = u_t \odot \tilde{h}_t + (1 - u_t) \odot h_{t-1}$$

$$\tilde{h} = \tanh(W[x_t] + U(r_t \odot h_{t-1}) + b)$$

$$u_t = \sigma(W_u[x_t] + U_u h_{t-1} + b_u)$$

$$r_t = \sigma(W_r[x_t] + U_r h_{t-1} + b_r)$$

### Long Short-Term Memory

[Hochreiter & Schmidhuber, NC1999;  
Gers, Thesis2001]

$$h_t = o_t \odot \tanh(c_t)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

$$\tilde{c}_t = \tanh(W_c[x_t] + U_c h_{t-1} + b_c)$$

$$o_t = \sigma(W_o[x_t] + U_o h_{t-1} + b_o)$$

$$i_t = \sigma(W_i[x_t] + U_i h_{t-1} + b_i)$$

$$f_t = \sigma(W_f[x_t] + U_f h_{t-1} + b_f)$$

Although computationally simpler, GRUs have similar performance to LSTMs  
Chung, Junyoung, et al. "Empirical evaluation of gated recurrent neural networks  
on sequence modeling." arXiv preprint arXiv:1412.3555 (2014)

[http://web.stanford.edu/class/cs20si/lectures/slides\\_11.pdf](http://web.stanford.edu/class/cs20si/lectures/slides_11.pdf)

© 2018 Ray Ptucha, Rochester Institute of Technology

40

## Learning Shakespeare

- LSTMs can learn structure and style in the data.
- Karpathy downloaded all the works of Shakespeare and concatenated them into a single (4.4MB) file.
- Train a 3-layer LSTM with 512 hidden nodes on each layer.
- After we train the network for a few hours Karpathy obtained samples such as:

© 2018 Ray Ptucha, Rochester Institute of Technology

41

PANDARUS:

Alas, I think he shall be come approached and the day  
When little strain would be attain'd into being never fed,  
And who is but a chain and subjects of his death,  
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,  
Breaking and strongly should be buried, when I perish  
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and  
my fair nudes begun out of the fact, to be conveyed,  
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

© 2018 Ray Ptucha, Rochester Institute of Technology

42

## Learning LaTeX

- The results above suggest that the model is actually quite good at learning complex syntactic structures.
- Karpathy and Johnston downloaded the raw Latex source file (a 16MB file) of a book on algebraic stacks/geometry and trained a multilayer LSTM.
- Amazingly, the resulting sampled LaTeX *almost* compiled.
- They had to step in and fix a few issues manually but then they get plausible looking math:

© 2018 Ray Ptucha, Rochester Institute of Technology

43

For  $\bigoplus_{i=1}^m \mathcal{L}_{m_i} = 0$ , hence we can find a closed subset  $\mathcal{H}$  in  $\mathcal{H}$  and any sets  $\mathcal{F}$  on  $X$ ,  $U$  is a closed immersion of  $S$ , then  $U \rightarrow T$  is a separated algebraic space.

*Proof.* Proof of (1). It also start we get

$$S = \text{Spec}(R) = U \times_X U \times_X U$$

and the comparicoly in the fibre product covering we have to prove the lemma generated by  $\coprod Z \times_U U \rightarrow V$ . Consider the maps  $M$  along the set of points  $\text{Sch}_{\text{fppf}}$  and  $U \rightarrow U$  is the fibre category of  $S$  in  $U$  in Section, ?? and the fact that any  $U$  affine, see Morphisms, Lemma ??, Hence we obtain a scheme  $S$  and any open subset  $W \subset U$  in  $\text{Sh}(G)$  such that  $\text{Spec}(R') \rightarrow S$  is smooth or an

$$U = \bigcup U_i \times_{S_i} U_i$$

which has a nonzero morphism we may assume that  $f_i$  is of finite presentation over  $S$ . We claim that  $\mathcal{O}_{X,x}$  is a scheme where  $x, x', s'' \in S'$  such that  $\mathcal{O}_{X,x'} \rightarrow \mathcal{O}_{X',s''}$  is separated. By Algebra, Lemma ?? we can define a map of complexes  $\text{GL}_{S'}(x'/S'')$  and we win.  $\square$

To prove study we see that  $\mathcal{F}|_U$  is a covering of  $\mathcal{X}'$ , and  $\mathcal{T}_i$  is an object of  $\mathcal{F}_{X/S}$  for  $i > 0$  and  $\mathcal{F}_i$  exists and let  $\mathcal{F}_i$  be a presheaf of  $\mathcal{O}_X$ -modules on  $\mathcal{C}$  as a  $\mathcal{F}$ -module. In particular  $\mathcal{F} = U/\mathcal{F}$  we have to show that

$$\tilde{M}^* = I^* \otimes_{\text{Spec}(k)} \mathcal{O}_{S,s} - i_X^{-1} \mathcal{F}$$

is a unique morphism of algebraic stacks. Note that

$$\text{Arrows} = (\text{Sch}/S)_{\text{fppf}}^{\text{opp}}, (\text{Sch}/S)_{\text{fppf}}$$

and


$$V = \Gamma(S, \mathcal{O}) \rightarrow (U, \text{Spec}(A))$$

is an open subset of  $X$ . Thus  $U$  is affine. This is a continuous map of  $X$  is the inverse, the groupoid scheme  $S$ .

*Proof.* See discussion of sheaves of sets.  $\square$

The result for prove any open covering follows from the less of Example ??, It may replace  $S$  by  $X_{\text{space}, \text{state}}$  which gives an open subspace of  $X$  and  $T$  equal to  $S_{\text{Zar}}$ , see Descent, Lemma ??, Namely, by Lemma ?? we see that  $R$  is geometrically regular over  $S$ .

This since  $\mathcal{F} \in \mathcal{F}$  and  $x \in \mathcal{G}$  the diagram



is a limit. Then  $\mathcal{G}$  is a finite type and assume  $S$  is a flat and  $\mathcal{F}$  and  $\mathcal{G}$  is a finite type  $\mathcal{F}$ . This is of finite type diagrams, and

- the composition of  $\mathcal{G}$  is a regular sequence,
- $\mathcal{O}_{X'}$  is a sheaf of rings.

$\square$

*Proof.* We have see that  $X = \text{Spec}(R)$  and  $\mathcal{F}$  is a finite type representable by algebraic space. The property  $\mathcal{F}$  is a finite morphism of algebraic stacks. Then the cohomology of  $X$  is an open neighbourhood of  $U$ .  $\square$

*Proof.* This is clear that  $\mathcal{G}$  is a finite presentation, see Lemmas ??, A reduced above we conclude that  $U$  is an open covering of  $\mathcal{C}$ . The functor  $\mathcal{F}$  is a "field"

$\mathcal{O}_{X,x} \rightarrow \mathcal{F}_x \rightarrow \mathcal{O}_{X_{\text{state}}} \rightarrow \mathcal{O}_{X,x}^* \mathcal{O}_{X,x}(\mathcal{O}_{X,x}^*)$  is an isomorphism of covering of  $\mathcal{O}_{X'}$ . If  $\mathcal{F}$  is the unique element of  $\mathcal{F}$  such that  $X$  is an isomorphism.

The property  $\mathcal{F}$  is a disjoint union of Proposition ?? and we can filtered set of presentations of a scheme  $\mathcal{O}_X$ -algebra with  $\mathcal{F}$  are opens of finite type over  $S$ . If  $\mathcal{F}$  is a scheme theoretic image points.  $\square$

If  $\mathcal{F}$  is a finite direct sum  $\mathcal{O}_{X'}$  is a closed immersion, see Lemma ??, This is a sequence of  $\mathcal{F}$  is a similar morphism.

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

© 2018 Ray Ptucha, Rochester Institute of Technology

44

## Learning C Code

The screenshot shows the GitHub repository for the Linux kernel source tree. At the top, there's a search bar and navigation links like 'Explore', 'Gist', 'Blog', and 'Help'. The repository name 'torvalds / linux' is prominently displayed. Below it, statistics show 520,037 commits, 1 branch, 420 releases, and 5,039 contributors. A list of recent merge commits is shown, including 'Merge branch 'drm-fixes' of git://people.freedesktop.org/~airlied/linux' and 'Merge git://git.kernel.org/pub/scm/linux/kernel/git/hab/target-pending'. On the right side, there are buttons for 'Watch', 'Star', 'Fork', and 'Pull requests', along with the 'HTTPS clone URL' and 'Clone in Desktop' options.

cs231n, lecture 10

© 2018 Ray Ptucha, Rochester Institute of Technology

45

```

static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << i))
            pipe = (in_use & UMXTHREAD_UNCCA) +
                ((count & 0x00000000ffffff8) & 0x000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &soffset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(m, "policy ");
}

```

cs231n, lecture 10

46

© 2018 Ray Ptucha, Rochester Institute of Technology

```

/*
 * Copyright (c) 2006-2010, Intel Mobile Communications. All rights reserved.
 *
 * This program is free software; you can redistribute it and/or modify it
 * under the terms of the GNU General Public License version 2 as published by
 * the Free Software Foundation.
 *
 * This program is distributed in the hope that it will be useful,
 * but WITHOUT ANY WARRANTY; without even the implied warranty of
 * MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
 * GNU General Public License for more details.
 *
 * You should have received a copy of the GNU General Public License
 * along with this program; if not, write to the Free Software Foundation,
 * Inc., 675 Mass Ave, Cambridge, MA 02139, USA.
 */

#include <linux/kexec.h>
#include <linux/errno.h>
#include <linux/io.h>
#include <linux/platform_device.h>
#include <linux/multi.h>
#include <linux/ckevent.h>

#include <asm/io.h>
#include <asm/prom.h>
#include <asm/e820.h>
#include <asm/system_info.h>
#include <asm/setew.h>
#include <asm/pgproto.h>

```

cs231n, lecture 10

47

© 2018 Ray Ptucha, Rochester Institute of Technology

```

#include <asm/io.h>
#include <asm/prom.h>
#include <asm/e820.h>
#include <asm/system_info.h>
#include <asm/setew.h>
#include <asm/pgproto.h>

#define REG_PG    vesa_slot_addr_pack
#define PFM_NOCOMP AFSR(0, load)
#define STACK_DDR(type)    (func)

#define SWAP_ALLOCATE(nr)    (e)
#define emulate_sigs()    arch_get_unaligned_child()
#define access_rw(TST)    asm volatile("movd %%esp, %0, %3" : : "r" (0)); \
    if (__type & DO_READ)

static void stat_PC_SEC __read_mostly offsetof(struct seq_argsqueue, \
    pC>[1]);

static void
os_prefix(unsigned long sys)
{
#ifdef CONFIG_PREEMPT
    PUT_PARAM_RAID(2, sel) = get_state_state();
    set_pid_sum((unsigned long)state, current_state_str(),
        (unsigned long)-1->lr_full; low;
}

```

cs231n, lecture 10

48

© 2018 Ray Ptucha, Rochester Institute of Technology

## The evolution of samples while training

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Let's examine how sampled text evolves while the LSTM model trains.
- For example, Karpathy trained an LSTM of Leo Tolstoy's War and Peace and then generated samples every 100 iterations of training.
- At iteration 100 the model samples random jumbles:

```

tyntd-iafhatawiaoihrdemot lytdws e ,tfti, astai f ogoh eoase rrranbyne 'nhthnee e
plia tklrge t o idoe ns,smtt h ne etie h,hregtrs nigtike,aoaenns lng

```

- The model is starting to get an idea about words separated by spaces- except sometimes it inserts two spaces. It also doesn't know that comma is almost always followed by a space.

© 2018 Ray Ptucha, Rochester Institute of Technology

49

## The evolution of samples while training

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- At 300 iterations we see that the model starts to get an idea about quotes and periods:

```
"Tmont thithey" fomesscerliund  
Keushey. Thom here  
sheulke, anmerenith ol sivh I lalterthend Bleipile shuw y fil on aseterlome  
coaniogennc Phe lism thond hon at. MeiDimorotion in ther thize."
```

- The words are now also separated with spaces and the model starts to get the idea about periods at the end of a sentence.
- At iteration 500:

```
we counter. He stutn co des. His stanted out one ofter that concoissions and was  
to gearang reay Jotrets and with fre colt of paitt thin wall. Which das stimn
```

© 2018 Ray Ptucha, Rochester Institute of Technology

50

## The evolution of samples while training

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- The model has now learned to spell the shortest and most common words such as "we", "He", "His", "Which", "and", etc. At iteration 700 we're starting to see more and more English-like text emerge:

```
Aftair fall unsuch that the hall for Prince Velzonski's that me of  
her hearly, and behs to so arwage fiving were to it beloge, pavu say falling misfort  
how, and Gogition is so overelical and ofter.
```

- At iteration 1200 we're now seeing use of quotations and question/exclamation marks. Longer words have now been learned as well:

```
"Kite vouch!" he repeated by her  
door. "But I would be done and quarts, feeling, then, son is people...."
```

© 2018 Ray Ptucha, Rochester Institute of Technology

51



## The evolution of samples while training

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Properly spelled words, quotations, names, and so on occur by about iteration 2000:


"Why do what that day," replied Natasha, and wishing to himself the fact the princess, Princess Mary was easier, fed in had oftened him.  
Pierre aking his soul came to the packs and drove up his father-in-law women.

- The model first discovers the general word-space structure and then rapidly starts to learn the words; First starting with the short words and then eventually the longer ones. Topics and themes that span multiple words (and in general longer-term dependencies) start to emerge only much later.

© 2018 Ray Ptucha, Rochester Institute of Technology

52


## Recurrent Networks for Character Prediction



© 2018 Ray Ptucha, Rochester Institute of Technology

53

## Recurrent Networks for Word Prediction



54

## Word2vec


- In the simplest form, we can start with a one-hot encoded vector of all words, and then learn a model which converts to a lower dimensional representation.
- Word2vec, glove, and skip-gram are popular metrics which encode words to a latent vector representation (~300 dimensions).
- Now we have a way to represent images, characters, and words as vectors.

© 2018 Ray Ptucha, Rochester Institute of Technology

55

# Sent2vec

- In the English to French translation, we have:



...but wait, this point in the RNN is a representation (sent2vec) of all the words in the English sentence!


- Now we have a way to represent images, characters, words, and sentences as vectors...can extend to paragraphs and documents...

© 2018 Ray Ptucha, Rochester Institute of Technology

56

# Image Captioning

RNN takes in a latent representation of an image, and generates a sequence.

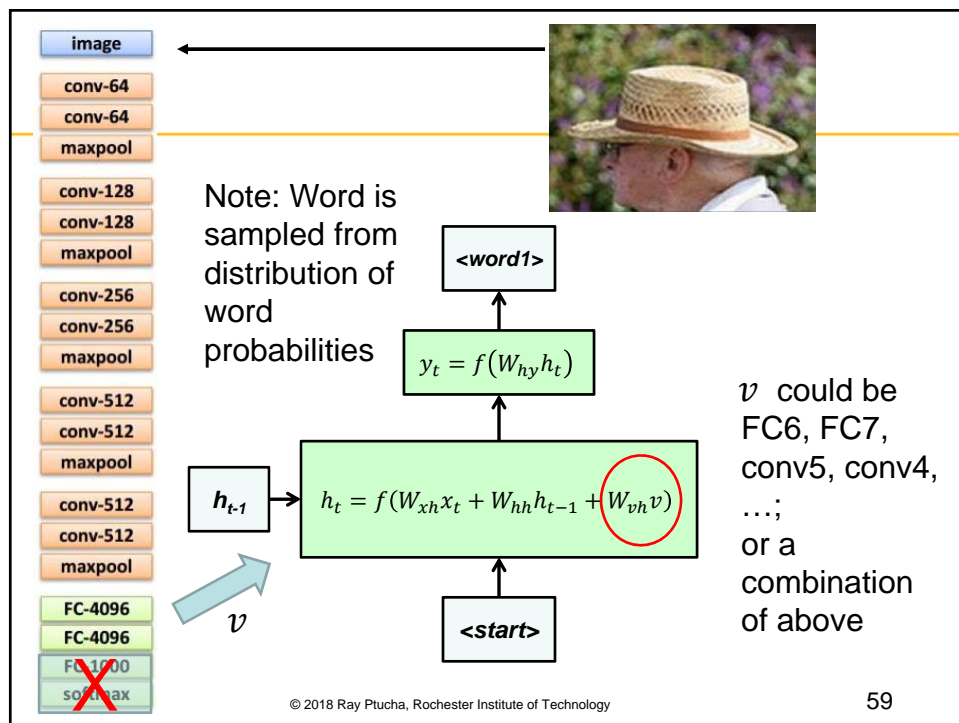
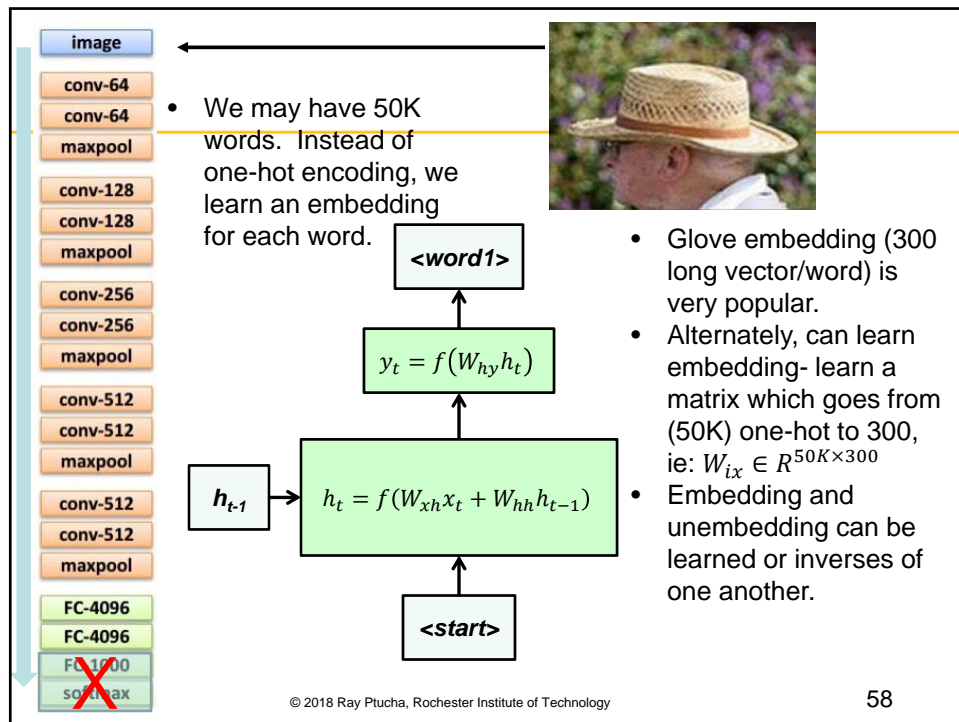


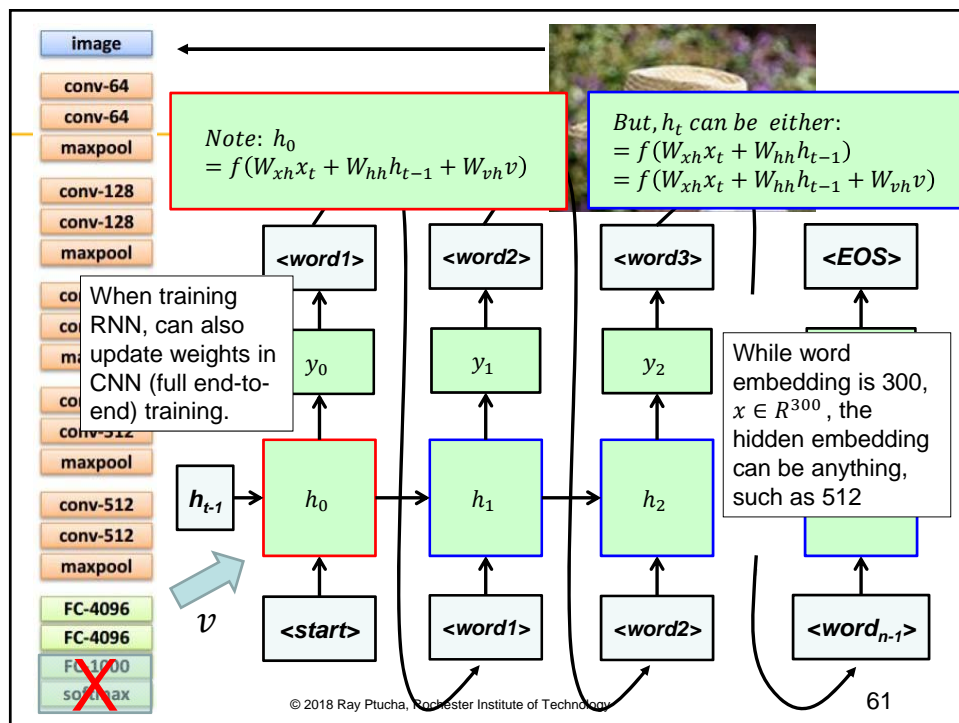
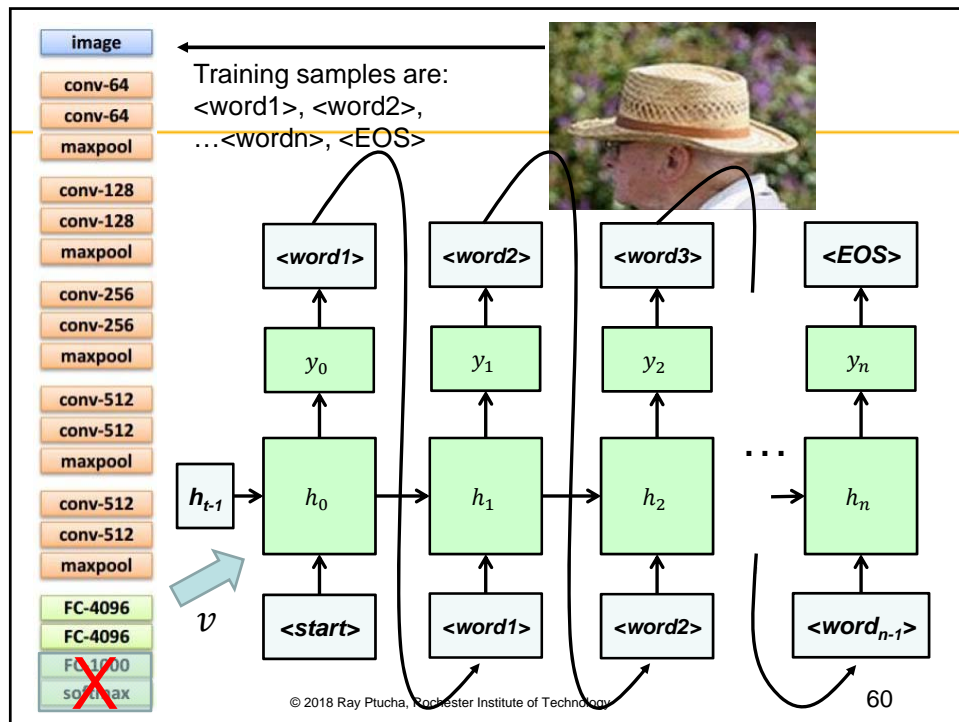
CNN helps represent an image as a numeric value. (image2vec)

Karpathy & Li, CVPR'15

© 2018 Ray Ptucha, Rochester Institute of Technology

57





## Data for Captioning

- Flickr8K
  - 8,000 images, from Flickr website, each with five captions
  - <http://nlp.cs.illinois.edu/HockenmaierGroup/8k-pictures.html>
- Flickr30K
  - 31,783 images, from Flickr website, each with five captions
  - <http://shannon.cs.illinois.edu/DenotationGraph/>
- MSCOCO
  - 80,000 training images, each with five captions
  - <http://mscoco.org/>

© 2018 Ray Ptucha, Rochester Institute of Technology

62

## Captioning Datasets

Amazon  
mechanical  
turkers do all  
labeling  
<https://www.mturk.com/mturk/welcome>




a man riding a bike on a dirt path through a forest.  
bicyclist raises his fist as he rides on desert dirt trail.  
this dirt bike rider is smiling and raising his fist in triumph.  
a man riding a bicycle while pumping his fist in the air.  
a mountain biker pumps his fist in celebration.

© 2018 Ray Ptucha, Rochester Institute of Technology

63

# MSCOCO Dataset



- A pile of wooden boxes filled with fruits and vegetables.
- An assortment of fruit in buckets for sale in a shop.
- An outdoor fruit stand with various types of fruits for sale.
- A display of crates of fruit on a city street.
- There are many crates with fruit and vegetables.

- A cat stands on a counter while a dog stands on the floor.
- A cat on the kitchen counter is looking down at a dog.
- A cat is looking at a dog rummage in the garbage.
- A cat on the counter and a dog on the ground in the kitchen.
- A cat stalking a dog on the kitchen floor.

© 2018 Ray Ptucha, Rochester Institute of Technology

64



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."



"boy is doing backflip on wakeboard."



"a young boy is holding a baseball bat."



"a cat is sitting on a couch with a remote control."



"a woman holding a teddy bear in front of a mirror."



"a horse is standing in the middle of a road."

© 2018 Ray Ptucha, Rochester Institute of Technology

Karpathy'15

65



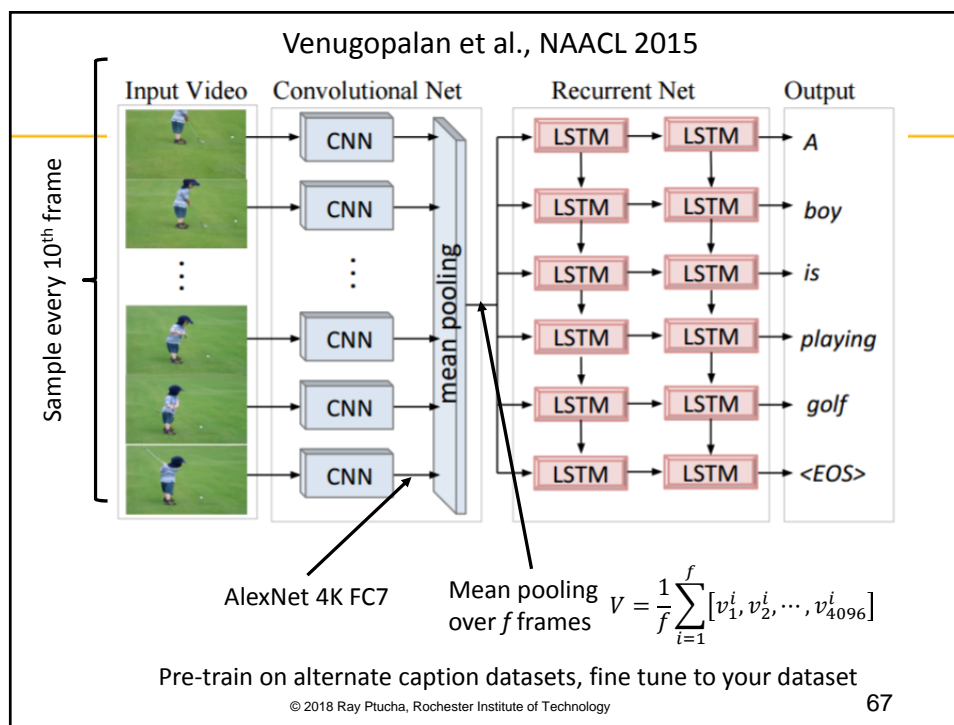
# Video Data for Captioning

MSVD: Microsoft Video Description Dataset  
MSR-VTT: Microsoft Research -Video to Text  
M-VAD: Movie description dataset M-VAD

	MSVD	MSR-VTT	M-VAD
#sentences	80,827	200,000	54,997
#sent. per video	~42	20	~1-2
vocab. size	9,729	24,282	16,307
avg. length	10.2s	14.8s	5.8s
#train video	1,200	6,513	36,921
#val. video	100	497	4,651
#test video	670	2,990	4,951


© 2018 Ray Ptucha, Rochester Institute of Technology

66





# Video Captioning



SV2T, Venugopalan, 2015

- Single LSTM for both encode and decode state.
- Two layer LSTM, 1000 hidden units each:
  - First LSTM learns video concepts
  - Second LSTM concentrates on language details.

© 2018 Ray Ptucha, Rochester Institute of Technology

68

## Thank you!!

Ray Ptucha

[rwpeec@rit.edu](mailto:rwpeec@rit.edu)



<https://www.rit.edu/mil>

© 2018 Ray Ptucha, Rochester Institute of Technology

69